

Classification Algorithms

Anirudh Singh, 2019B5A70948H
Bharath Variar, 2019B5A70930H

Assignment-1 Report
BITS F266: Machine Learning
Department of Computer Science and Information
BITS Pilani, Hyderabad Campus
March 2023

Contents

1	Perceptron Learning Algorithm	2
1.1	Data Cleaning	2
1.2	Perceptron Algorithm	2
1.3	Task 1	2
1.4	Task 2	4
1.5	Task 3	5
1.6	Results	6
2	Fisher's Linear Discriminant Analysis	6
2.1	Task 1	6
2.2	Task 2	7
3	Logistic Regression	8
3.1	Implementation	8
3.1.1	Vanilla Gradient Descent	8
3.1.2	Stochastic Gradient Descent	10
3.1.3	MiniBatch Gradient Descent	12
3.2	Task 1: Raw Data	13
3.3	Task 2: Normalised Data	13
	References	16

1 Perceptron Learning Algorithm

1.1 Data Cleaning

For the perceptron algorithm to work efficiently, we first remove the `id` column since it consists of large values that can disturb the learning process while not contributing to it semantically, as the `id` assigned to any sample does not determine the type of the tumour. We then map the target attribute, `diagnosis`, to numeric values using the `map` function. We assign value 1 to Malignant and -1 to Benign tumours.

1.2 Perceptron Algorithm

The Perceptron algorithm is implemented using a `Perceptron` class which is implemented as follows:

```
1 class Perceptron:
2     def __init__(self, lr=0.001, epochs=10):
3         self.lr = lr
4         self.epochs = epochs
5         self.weights = None
6
7     def train(self, X, y):
8         n_samples, n_features = X.shape
9         self.weights = np.zeros(n_features)
10
11         for _ in range(self.epochs):
12             y_h=np.dot(X,self.weights)
13             for idx, x_i in X.iterrows():
14                 y_hat = np.dot(x_i, self.weights)
15                 if y_hat*y[idx]<=0:
16                     self.weights += x_i*y[idx]
17
18     def predict(self, X):
19         y_hat=np.dot(X, self.weights)
20         return np.where(y_hat >= 0, 1, -1)
```

1.3 Task 1

In this task, we train the perceptron algorithm on the raw data (PM1) and on shuffled data (PM2). We can see that PM1 gives an accuracy of 72.3% while sPM2 gives an accuracy of 79.3%, suggesting that the data is linearly separable and shuffling the data improves the learning of the perceptron algorithm.

Task 1: Perceptron on raw data (PM1)

```
In [6]: 1 y = df['diagnosis']
2 X = df.drop(['diagnosis'], axis=1)
3 split_idx = int(len(df) * 0.67)
4 X_train, X_test = X[:split_idx], X[split_idx:]
5 y_train, y_test = y[:split_idx], y[split_idx:]
```

```
In [7]: 1 PM1 = Perceptron(lr = 0.01, epochs = 500)
2 PM1.train(X_train, y_train)
3 y_pred1 = PM1.predict(X_test)
4
5 accuracy1 = evaluate(y_test, y_pred1)
```

Confusion Matrix: {'true_positive': 42, 'true_negative': 94, 'false_positive': 51, 'false_negative': 1}
Accuracy: 72.34042553191489%
Precision: 45.16129032258065%
Recall: 97.67441860465117%

Figure 1: PM1

Task 1: Perceptron on shuffled data (PM2)

```
In [8]: 1 df_shuffled = df.sample(frac=1, random_state=42) #shuffling dataset
2 y_shuffled = df_shuffled['diagnosis']
3 X_shuffled = df_shuffled.drop(['diagnosis'], axis=1)
4 split_idx = int(len(df_shuffled) * 0.67)
5 X_train, X_test = X_shuffled[:split_idx], X_shuffled[split_idx:]
6 y_train, y_test = y_shuffled[:split_idx], y_shuffled[split_idx:]
```

```
In [9]: 1 PM2 = Perceptron(lr = 0.01, epochs = 500)
2 PM2.train(X_train, y_train)
3 y_pred2 = PM1.predict(X_test)
4
5 accuracy2 = evaluate(y_test, y_pred2)
```

Confusion Matrix: {'true_positive': 73, 'true_negative': 76, 'false_positive': 36, 'false_negative': 3}
Accuracy: 79.25531914893617%
Precision: 66.97247706422019%
Recall: 96.05263157894737%

Figure 2: PM2

1.4 Task 2

In this task, we attempt to improve the learning by normalising the data based on the formula:

$$X = \frac{X - \mu}{\sigma} \quad (1)$$

where μ, σ are the mean and standard deviation of X respectively.

We call the model that trains on normalised data PM3, and it classifies with an accuracy of 93.1% as shown below. Comparing this model with PM1, we see that normalising the data greatly improves the accuracy of classification, and this is because the change in weights in each iteration is not large since the data is small due to normalisation.

Task 2: Perceptron on normalised data without shuffling (PM3)

Normalising

```
In [10]: 1 mean = np.mean(X, axis=0)
          2 stddev = np.std(X, axis=0)
          3 X_normalised = (X - mean) / stddev
          4 split_idx = int(len(df) * 0.67)
          5 X_train, X_test = X_normalised[:split_idx], X_normalised[split_idx:]
          6 y_train, y_test = y[:split_idx], y[split_idx:]

In [11]: 1 X_normalised.head()
```

Out[11]:

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	1.094615	-2.073335	1.264830	0.984375	1.568466	3.283515
1	1.818905	-0.353632	1.678196	1.908708	-0.826962	-0.487072
2	1.571860	0.456187	1.559507	1.558884	0.942210	1.052926
3	-0.749801	0.253732	-0.585900	-0.764464	3.283553	3.402909
4	1.740300	-1.151816	1.768236	1.826229	0.280372	0.539340

5 rows × 30 columns

```
In [12]: 1 PM3 = Perceptron(lr = 0.01, epochs = 500)
          2 PM3.train(X_train, y_train)
          3 y_pred3 = PM3.predict(X_test)
          4
          5 accuracy3 = evaluate(y_test, y_pred3)
```

Confusion Matrix: {'true_positive': 40, 'true_negative': 135, 'false_positive': 10, 'false_negative': 3}
Accuracy: 93.08510638297872%
Precision: 80.0%
Recall: 93.02325581395348%

Figure 3: PM3

1.5 Task 3

In the final task for the perceptron algorithm, we train the model, PM4 on normalised data with the features shuffled. PM4 classifies with an accuracy of 93.1%, exactly the same as PM3, which implies that while shuffling data improves accuracy, shuffling the feature tuple does not. This still performs much better than PM1, but this is due to normalisation and not the shuffled feature tuple, since the accuracies of both PM3 and PM4 are identical.

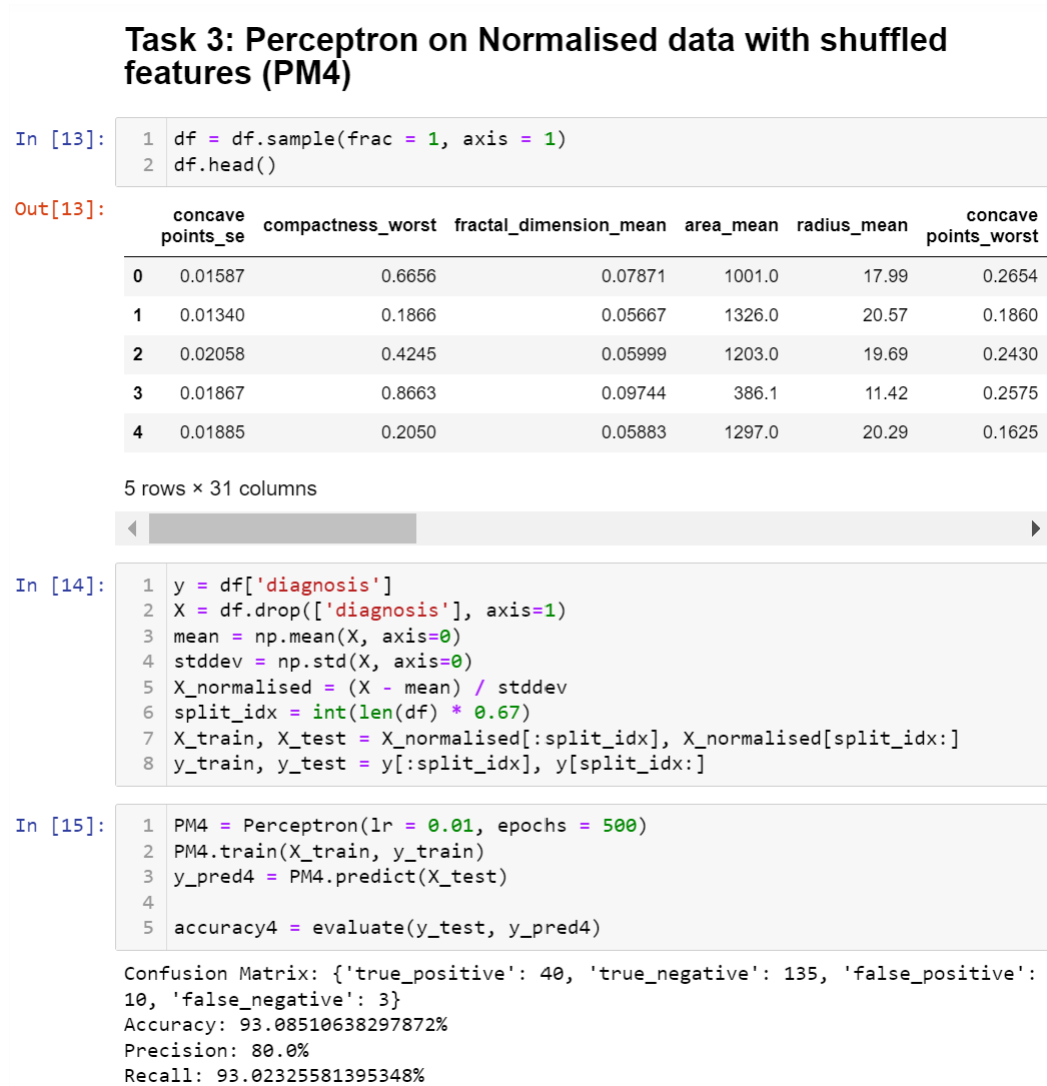


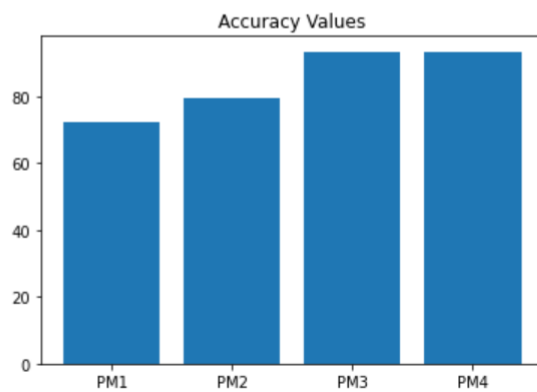
Figure 4: PM4

1.6 Results

We finally plot the accuracies of all the models in a bar chart for a comparative study as shown:

Results

```
In [16]: 1 accuracy = {'PM1': accuracy1,
2             'PM2': accuracy2,
3             'PM3': accuracy3,
4             'PM4': accuracy4}
5 names = ['PM1', 'PM2', 'PM3', 'PM4']
6 values = [accuracy1, accuracy2, accuracy3, accuracy4]
7 plt.bar(range(len(accuracy)), values, tick_label = names)
8 plt.title("Accuracy Values")
9 plt.show()
10 print (accuracy)
```



```
{'PM1': 72.34042553191489, 'PM2': 79.25531914893617, 'PM3': 93.08510638297872,
'PM4': 93.08510638297872}
```

Figure 5: Results

2 Fisher's Linear Discriminant Analysis

2.1 Task 1

We use the linear discriminant model from `scikitlearn` and call it FLDM1. Implementation is done as shown below:

Task 1: FLDM on raw data (FLDM1)

```
In [3]: 1 # build the FLDM model
2 fldm1 = LinearDiscriminantAnalysis(n_components=1)
3 fldm1.fit(X_train, y_train)
4
5 # project the training data onto the 1-dimensional FLDM space
6 X_train_lda = fldm1.transform(X_train)

In [4]: 1 # find the decision boundary in the 1-dimensional FLDM space
2 mean_pos = np.mean(X_train_lda[y_train == 1])
3 mean_neg = np.mean(X_train_lda[y_train == -1])
4 std_pos = np.std(X_train_lda[y_train == 1])
5 std_neg = np.std(X_train_lda[y_train == -1])
6
7 threshold = (mean_pos + mean_neg) / 2
8
9 # project the testing data onto the 1-dimensional FLDM space
10 X_test_lda = fldm1.transform(X_test)
11
12 # evaluate the performance of the model on the testing data
13 y_pred = np.where(X_test_lda > threshold, 1, -1)
14 accuracy_fldm1 = evaluate(y_test, y_pred)

Confusion Matrix: {'true_positive': 41, 'true_negative': 142, 'false_positive':
3, 'false_negative': 2}
Accuracy: 97.34042553191489%
Precision: 93.18181818181819%
Recall: 95.34883720930233%
```

Figure 6: FLDM1

We obtain an accuracy of 97.3%.

2.2 Task 2

In this task, we shuffle the feature tuple and then call this model FLDM2. The implementation is shown below. As we can see, the accuracy is 97.3%, identical to that produced by FLDM1, and this implies that shuffling the features does not change the learning for this algorithm.

Task 2: FLDM on data with shuffled columns (FLDM2)

```
In [8]: 1 # randomly shuffle the order of features in the training data
2 np.random.seed(42)
3 n_features = X_train.shape[1]
4 feature_order = np.random.permutation(n_features)
5 X_train_shuffled = X_train[X_train.columns[feature_order]]
6 X_train_shuffled.head()
```

```
Out[8]:
```

	concave points_worst	compactness_se	area_worst	concave points_se	symmetry_mean	fractal_dimension_mean	symmetry_worst	smoothness_worst	perimeter_se	radius_m
0	0.2654	0.04904	2019.0	0.01587	0.2419	0.07871	0.4601	0.1622	8.589	1
1	0.1860	0.01308	1956.0	0.01340	0.1812	0.05667	0.2750	0.1238	3.398	2
2	0.2430	0.04006	1709.0	0.02058	0.2069	0.05999	0.3613	0.1444	4.585	1
3	0.2575	0.07458	567.7	0.01867	0.2597	0.09744	0.6638	0.2098	3.445	1
4	0.1625	0.02461	1575.0	0.01885	0.1809	0.05883	0.2364	0.1374	5.438	2

5 rows x 30 columns

```
In [6]: 1 # build the FLDM model
2 fldm2 = LinearDiscriminantAnalysis(n_components=1)
3 fldm2.fit(X_train_shuffled, y_train)
4
5 # project the training data onto the 1-dimensional FLDM space
6 X_train_lda = fldm2.transform(X_train_shuffled)
```

```
In [7]: 1 # find the decision boundary in the 1-dimensional FLDM space
2 mean_pos = np.mean(X_train_lda[y_train == 1])
3 mean_neg = np.mean(X_train_lda[y_train == -1])
4 std_pos = np.std(X_train_lda[y_train == 1])
5 std_neg = np.std(X_train_lda[y_train == -1])
6
7 threshold = (mean_pos + mean_neg) / 2
8
9 # shuffle the order of features in the testing data
10 X_test_shuffled = X_test[X_test.columns[feature_order]]
11
12 # project the testing data onto the 1-dimensional FLDM space
13 X_test_lda = fldm2.transform(X_test_shuffled)
14
15 # evaluate the performance of the model on the testing data
16 y_pred = np.where(X_test_lda > threshold, 1, -1)
17 accuracy2 = evaluate(y_test, y_pred)
```

Confusion Matrix: {'true_positive': 41, 'true_negative': 142, 'false_positive': 3, 'false_negative': 2}
Accuracy: 97.34042553191489%
Precision: 93.18181818181819%
Recall: 95.34883720930233%

Figure 7: FLDM2

3 Logistic Regression

3.1 Implementation

We have implemented vanilla, stochastic and minibatch gradient descent in the object-oriented paradigm by defining various classes as shown below:

3.1.1 Vanilla Gradient Descent

LogisticRegressionGD as follows:

```
1 class LogisticRegressionGD:
```

```

2 def __init__(self, learning_rate=0.01, n_iters=1000, random_state=
  None, threshold=0.5):
3     self.learning_rate = learning_rate
4     self.n_iters = n_iters
5     self.random_state = random_state
6     self.threshold=threshold
7
8     def sigmoid(self, x):
9         return 1 / (1 + np.exp(-x))
10
11    def fit(self, X, y):
12        n_samples, n_features = X.shape
13        self.weights = np.zeros(n_features)
14        self.bias = 0
15        self.costs = []
16
17        # set random seed for reproducibility
18        if self.random_state is not None:
19            np.random.seed(self.random_state)
20
21        # gradient descent
22        for i in range(self.n_iters):
23            # calculate predicted probabilities and gradients
24            linear_model = np.dot(X, self.weights) + self.bias
25            y_pred = self.sigmoid(linear_model)
26            dw = np.dot(X.T, (y_pred - y)) / n_samples
27            db = np.sum(y_pred - y) / n_samples
28
29            # update weights and bias
30            self.weights -= self.learning_rate * dw
31            self.bias -= self.learning_rate * db
32
33            # calculate cost and add to list for graphing
34            y_pred = self.sigmoid(np.dot(X, self.weights) + self.
bias)
35            y_pred[y_pred == 0] = 1e-15 # add small constant value
to avoid NaN in cost
36            y_pred[y_pred == 1] = 1 - 1e-15 # add small constant
value to avoid NaN in cost
37            cost = -1/n_samples * np.sum(y * np.log(y_pred) + (1-y)
* np.log(1-y_pred))
38            self.costs.append(cost)
39
40    def predict(self, X):
41        linear_model = np.dot(X, self.weights) + self.bias
42        y_pred = self.sigmoid(linear_model)
43        y_pred_class = [1 if i > self.threshold else 0 for i in
y_pred]

```

```

44         return y_pred_class
45
46     def plot_cost(self):
47         fig, ax = plt.subplots(figsize=(10, 8))
48         plt.plot(np.arange(1, len(self.costs)+1), self.costs)
49         plt.xlabel('Iterations')
50         plt.ylabel('Cost')
51         plt.title('Gradient Descent Cost Graph')
52         plt.show()
53         fig.savefig('unnormalizedLR_graphs/'+str(len(str(self.
learning_rate)))+str(self.threshold)[-1]+".png")

```

3.1.2 Stochastic Gradient Descent

```

1 class LogisticRegressionSGD:
2     def __init__(self, learning_rate=0.01, n_iters=1000, batch_size
=1, random_state=None, threshold=0.5):
3         self.learning_rate = learning_rate
4         self.n_iters = n_iters
5         self.batch_size = batch_size
6         self.random_state = random_state
7         self.threshold=threshold
8
9     def sigmoid(self, x):
10         return 1 / (1 + np.exp(-x))
11
12     def fit(self, X, y):
13         n_samples, n_features = X.shape
14         self.weights = np.zeros(n_features)
15         self.bias = 0
16         self.costs = []
17
18         # set random seed for reproducibility
19         if self.random_state is not None:
20             np.random.seed(self.random_state)
21
22         # stochastic gradient descent
23         for i in range(self.n_iters):
24             # shuffle data
25             idx = np.arange(n_samples)
26             np.random.shuffle(idx)
27             X_shuffled = X[idx]
28             y_shuffled = y[idx]
29
30             # loop over batches
31             for j in range(0, n_samples, self.batch_size):
32                 # get mini-batch

```

```

33         X_batch = X_shuffled[j:j+self.batch_size]
34         y_batch = y_shuffled[j:j+self.batch_size]
35
36         # calculate predicted probabilities and gradients
37         linear_model = np.dot(X_batch, self.weights) + self.
bias
38         y_pred = self.sigmoid(linear_model)
39         dw = np.dot(X_batch.T, (y_pred - y_batch)) / self.
batch_size
40         db = np.sum(y_pred - y_batch) / self.batch_size
41
42         # update weights and bias
43         self.weights -= self.learning_rate * dw
44         self.bias -= self.learning_rate * db
45
46         # calculate cost and add to list for graphing
47         y_pred = self.sigmoid(np.dot(X, self.weights) + self.
bias)
48         y_pred[y_pred == 0] = 1e-15 # add small constant value
to avoid NaN in cost
49         y_pred[y_pred == 1] = 1 - 1e-15 # add small constant
value to avoid NaN in cost
50         cost = -1/n_samples * np.sum(y * np.log(y_pred) + (1-y)
* np.log(1-y_pred))
51         self.costs.append(cost)
52
53     def predict_proba(self, X):
54         linear_model = np.dot(X, self.weights) + self.bias
55         y_pred = self.sigmoid(linear_model)
56         return y_pred
57
58     def predict(self, X):
59         y_pred_proba = self.predict_proba(X)
60         y_pred_class = [1 if i > self.threshold else 0 for i in
y_pred_proba]
61         return y_pred_class
62
63     def plot_cost(self):
64         fig, ax = plt.subplots(figsize=(10, 8))
65         plt.plot(np.arange(1, len(self.costs)+1), self.costs)
66         plt.xlabel('Iterations')
67         plt.ylabel('Cost')
68         plt.title('Stochastic Gradient Descent Cost Graph')
69         plt.show()
70         fig.savefig('unnormalizedLR_graphs/'+ "SGD"+str(len(str(self.
learning_rate)))+str(self.threshold)[-1]+".png")

```

3.1.3 MiniBatch Gradient Descent

```
1 class LogisticRegressionMiniBatchGD:
2     def __init__(self, learning_rate=0.01, n_iters=1000, batch_size
3         =32, random_state=None, threshold=0.5):
4         self.learning_rate = learning_rate
5         self.n_iters = n_iters
6         self.batch_size = batch_size
7         self.random_state = random_state
8         self.threshold=threshold
9
10    def sigmoid(self, z):
11        return 1 / (1 + np.exp(-z))
12
13    def fit(self, X, y):
14        # initialize weights and bias to zero
15        self.weights = np.zeros(X.shape[1])
16        self.bias = 0
17
18        # initialize costs list for storing costs at each iteration
19        self.costs = []
20
21        # set random seed for reproducibility
22        n_samples, n_features = X.shape
23        if self.random_state is not None:
24            np.random.seed(self.random_state)
25
26        # minibatch gradient descent
27        for i in range(self.n_iters):
28            # shuffle data
29            idx = np.arange(X.shape[0])
30            np.random.shuffle(idx)
31            X = X[idx]
32            y = y[idx]
33
34            # loop over batches
35            for j in range(0, X.shape[0], self.batch_size):
36                # get minibatch
37                X_batch = X[j:j+self.batch_size]
38                y_batch = y[j:j+self.batch_size]
39
40                # calculate predicted probabilities and gradients
41                y_pred = self.sigmoid(np.dot(X_batch, self.weights)
42                    + self.bias)
43                dw = np.dot(X_batch.T, (y_pred - y_batch)) / self.
44                    batch_size
45                db = np.sum(y_pred - y_batch) / self.batch_size
```

```

44         # update weights and bias
45         self.weights -= self.learning_rate * dw
46         self.bias -= self.learning_rate * db
47
48         # calculate cost and add to list for graphing
49         y_pred = self.sigmoid(np.dot(X, self.weights) + self.
bias)
50         y_pred[y_pred == 0] = 1e-15 # add small constant value
to avoid NaN in cost
51         y_pred[y_pred == 1] = 1 - 1e-15 # add small constant
value to avoid NaN in cost
52         cost = -1/n_samples * np.sum(y * np.log(y_pred) + (1-y)
* np.log(1-y_pred))
53         self.costs.append(cost)
54
55         return self
56
57     def predict_proba(self, X):
58         return self.sigmoid(np.dot(X, self.weights) + self.bias)
59
60     def predict(self, X):
61         return np.where(self.predict_proba(X) > self.threshold, 1,
0)
62
63     def plot_cost(self):
64         fig, ax = plt.subplots(figsize=(10, 8))
65         plt.plot(range(1, len(self.costs) + 1), self.costs)
66         plt.xlabel('Iteration')
67         plt.ylabel('Cost')
68         plt.title('Logistic Regression Cost Graph')
69         plt.show()
70         fig.savefig('unnormalizedLR_graphs/'+ "minibatchGD"+str(len(
str(self.learning_rate)))+str(self.threshold)[-1]+".png")

```

3.2 Task 1: Raw Data

Varying thresholds did not change the accuracies much, but we got the best accuracy of 92.6% when the threshold was 0.5 and the learning rate of 0.001 using vanilla gradient descent, while we got the worst accuracy of 32.9% when the threshold was 0.6 and the learning rate 0.01.

3.3 Task 2: Normalised Data

In this task, we run the same algorithms after normalising data using the code shown below:

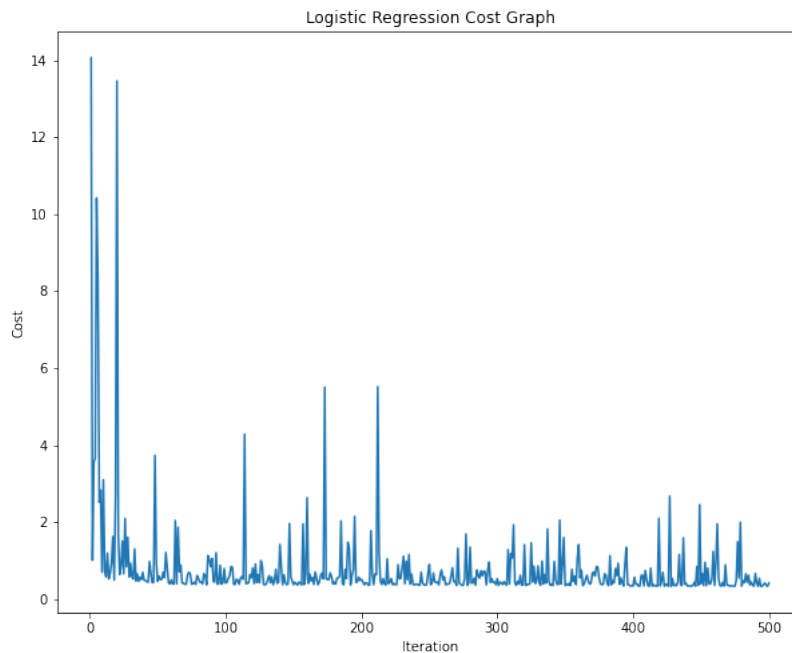


Figure 8: Mini-batch GD on raw Data

```

1  y = df["diagnosis"]
2  X = df.drop(["diagnosis"], axis=1)
3  mean = np.mean(X, axis=0)
4  stddev = np.std(X, axis=0)
5  X = (X - mean) / stddev
6  df = pd.concat([X, y], axis=1)

```

After normalising, while accuracy remained relatively small, from 92.6% to 98.9%, we noticed that the graphs for cost curves were significantly more stable, and this is because there are no large changes in weights in each iteration.

We obtained the highest accuracy of 98.9% with a learning rate of 0.001 and a threshold of 0.5 using the minibatch gradient descent. While the worst was an accuracy of 96.2% with a learning rate, of 0.01, and the threshold of 0.4 using stochastic gradient descent.

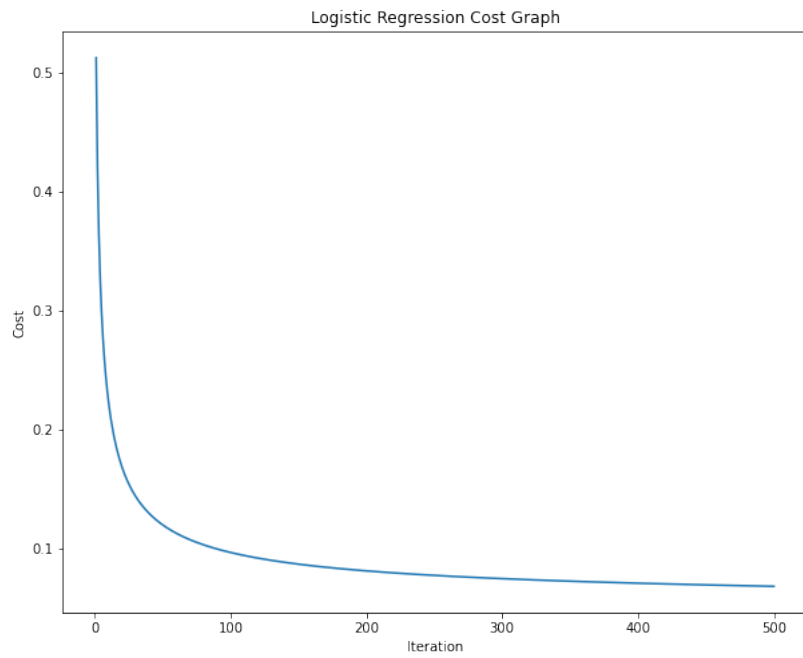


Figure 9: Mini-batch GD on normalised Data

Note: It, however, should be noted that with extremely low learning rates (0.0001), the gradient descent does not converge to the global minima in > 500 iterations.

References

- [1] [GitHub Repository](#)